

Retrieval at Scale: Results

Transformer models trained on the corpus reach 71.3 percent exact match on the held-out split. The gain over the recurrent baseline is largest on questions that need two or more hops of reasoning across documents.

Inference cost scales linearly with context length up to 128,000 tokens. Beyond that length the retrieval variant is cheaper by a factor of 3.2 while losing only 0.4 points of accuracy on the same split.